

Assignment 1: CS 763, Computer Vision

Due: 26th Jan before 11:00 pm

Remember the honor code while submitting this (and every other) assignment. All members of the group should work on and understand all parts of the assignment. We will adopt a zero-tolerance policy against any violation.

Submission instructions: You should ideally type out all the answers in Word (with the equation editor) or using LaTeX. In either case, prepare a pdf file. For assignment submission, follow the instructions for arrangement of folders and subfolders as given in https://github.com/cs763/Spring2018/blob/master/assignments/assignment1/HW1_CameraGeometry.tar.gz. Create a single zip or rar file obeying the aforementioned structure and name it as follows: A1-IdNumberOfFirstStudent-IdNumberOfSecondStudent-IdNumberOfThirdStudent.zip. (If you are doing the assignment alone, the name of the zip file is A1-IdNumber.zip). Upload the file on moodle BEFORE 11:00 pm on 26th January. Late assignments will be assessed a penalty of 50% per day late. Note that only one student per group should upload their work on moodle. Please preserve a copy of all your work until the end of the semester. If you have difficulties, please do not hesitate to seek help from me.

1. In this question, you will calibrate your own camera using a checkerboard object. Print two (or three) copies of the checkerboard image (checkbox.jpg) provided in the input folder and stick them to two (or three) orthogonal planes (wall corner). Click a picture of the checkerboard from your phone. Camera calibration requires 2D and 3D correspondences. Create a dataset that contains XYZ coordinates of N points marked out on the wall checkerboard and also the XY coordinates of the corresponding points on the image. Now, write a MATLAB program that estimates the 3×4 projection matrix P and then decompose it into the intrinsics and extrinsics. Your code should follow the following steps:

- Normalize the data such that the centroid of 2D and 3D points are at origin and the average Euclidean distance of 2D and 3D points from the origin is $\sqrt{2}$ and $\sqrt{3}$, respectively. Find the transformation matrices T and U that achieve this for 2D and 3D respectively, *i.e.*, $\hat{\mathbf{x}} = T\mathbf{x}$ and $\hat{\mathbf{X}} = U\mathbf{X}$ where \mathbf{x} and \mathbf{X} are the unnormalized 2D and 3D points in homogeneous coordinates.
- Estimate the normalized projection matrix $\hat{\mathbf{P}}$ using the DLT method as discussed in the class. Denormalize the projection matrix $\hat{\mathbf{P}}$ ($\mathbf{P} = T^{-1}\hat{\mathbf{P}}U$).
- Decompose the projection matrix $\mathbf{P} = K[R| -RX_0]$ into intrinsic matrix K , rotation matrix R and the camera center X_0 . K and R can be estimated using RQ decomposition.
- Verify that the projection matrix is correctly estimated by computing the RMSE between the 2D points marked by you and the estimated 2D projections of the marked 3D points. Visualize the points on the image and include them in the report. Also mention why it is a good idea to normalize the points before performing DLT. **[5+15+5+5 points]**

2. In this question, we will perform distortion correction. Consider the image rad_checkbox.jpg provided in the input folder. The image has been distorted using the following distortion model and distortion parameters $\mathbf{q} = \{q_1, q_2\}$:

$$\mathbf{x}_d = \mathbf{x}_u(1 + q_1r + q_2r^2)$$

where $q_1 = 1$ and $q_2 = 0.5$ and $r = \|\mathbf{x}_u\|_2$. First, normalize x_d between $(-1, 1)$. Next, iteratively find x_u using the update formula until convergence (should converge quickly):

$$\mathbf{x}_u \leftarrow \begin{bmatrix} 1 & 0 & \Delta x(\mathbf{x}_d, \mathbf{q}) \\ 0 & 1 & \Delta y(\mathbf{x}_d, \mathbf{q}) \\ 0 & 0 & 1 \end{bmatrix} \mathbf{x}_d$$

Finally, warp the image intensities at \mathbf{x}_d to \mathbf{x}_u using MATLAB's `interp2` function. Display the undistorted image and include the code in the submission. [20 points]

3. Given a projective transformation matrix:

$$M = \begin{bmatrix} 0 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 0 \end{bmatrix}$$

find the image of a hyperbola $y = 1/x$ under the transform M . (Represent the hyperbola in parametric form $(t, 1/t), t \neq 0$ in homogeneous coordinates and find the projection.) Comment on the shape of the image of hyperbola obtained. [10 points]

4. Here you will develop an application of the concept of vanishing points to camera calibration. Consider two images (I_P and I_Q) of a non-planar scene taken with two pinhole cameras having unknown focal lengths f_p and f_q respectively. Both cameras produce images on a Cartesian grid with aspect ratio of 1 and unknown resolution s_p and s_q respectively. The orientations and positions of the two cameras are related by an unknown rotation (given by a 3×3 rotation matrix \mathbf{R}) and an unknown translation (given by a 1×3 vector \mathbf{t}). Note that ‘position’ here refers to the location of the camera pinhole, and ‘orientation’ refers to the XYZ axes of the camera coordinate system. In both I_P and I_Q , suppose you accurately mark out the corresponding vanishing points of three mutually perpendicular directions ℓ_1, ℓ_2, ℓ_3 in the scene. (Obviously, all three directions were visible from both cameras). Let the vanishing points have coordinates (p_{1x}, p_{1y}) , (p_{2x}, p_{2y}) and (p_{3x}, p_{3y}) in I_P , and (q_{1x}, q_{1y}) , (q_{2x}, q_{2y}) and (q_{3x}, q_{3y}) in I_Q . Note that these coordinates are in terms of pixel units and the correspondences are known. Given all this information, can you infer \mathbf{R} ? Can you infer \mathbf{t} ? Can you infer f_p and f_q ? Can you infer s_p and s_q ? Explain how (or why not). (Hint: You can start off by assuming that you knew all the intrinsic parameters and work yourself upwards from there). [20 points]
5. a) Prove (algebraically) that the projections (in image plane) of any two parallel lines L_1, L_2 in \mathbb{R}^3 have an intersection point, the vanishing point.
b) Prove (algebraically) that the vanishing points corresponding to three (different) sets of parallel lines on a 3D plane are collinear in the image plane. [5 + 5 points]
6. **Image Forensics:** Consider the image painting.jpg provided in the input directory of Question 6. It is the image of the painting *Flagellation of Christ* by Piero della Francesca. Assume that the green line on the image is the horizon line and the also that the height of Christ, \mathbf{R} is 180cm. The painting has been lauded for its correct depiction of linear perspectives. Using your knowledge of Computer Vision, can you find the height, \mathbf{H} , of the person on the right? Mention the estimated height in the report along with a detailed description of how you arrived at the answer. [10 points]